

Dhruv Sreenivas

✉ dhruv.sreenivas@mila.quebec

☎ +1 (978) 424-7614

🌐 linkedin.com/in/dhruvsreenivas

🐙 github.com/dhruvsreenivas

EDUCATION

Mila, Université de Montréal

Ph.D. in Computer Science

GPA: 4.15/4.0

Sep 2024 - Present

Montréal, QC, Canada

Advisor: Pablo Samuel Castro

- Relevant courses: Representation Learning (PhD), Probabilistic Graphical Models

Cornell University

M.S. in Computer Science

GPA: 4.30/4.0

Aug 2021 - Aug 2023

Ithaca, NY, USA

Advisors: Wen Sun, Robert Kleinberg

- Teaching assistant for CS 2110, CS 4789, CS 6756 (PhD), CS 6789 (PhD)
- Relevant courses: Advanced Topics in Machine Learning (PhD), Deep Generative Models (PhD), Foundations of Reinforcement Learning (PhD), Machine Learning Theory (PhD), Machine Learning for Feedback Systems (PhD), Advanced Machine Learning Systems (PhD)

Cornell University

B.S. in Computer Science, Mathematics

GPA: 3.66/4.0

Aug 2018 - May 2021

Ithaca, NY, USA

PUBLICATIONS

(*: equal contribution)

1. Olya Mastikhina*, **Dhruv Sreenivas***, Pablo Samuel Castro
Optimistic Critics Can Empower Small Actors
Reinforcement Learning Conference (**RLC**), **2025**
2. Jonathan Chang, **Dhruv Sreenivas**, Yingbing Huang, Kianté Brantley, Wen Sun
Adversarial Imitation Learning via Boosting
International Conference on Learning Representations (**ICLR**), **2024**
3. Shachi Deshpande, Kaiwen Wang, **Dhruv Sreenivas**, Zheng Li, Volodymyr Kuleshov
Deep Multi-Modal Structural Equations for Causal Effect Estimation with Unstructured Proxies
Neural Information Processing Systems (**NeurIPS**), **2022**
4. Jonathan Chang, Masatoshi Uehara, **Dhruv Sreenivas**, Rahul Kidambi, Wen Sun
Mitigating Covariate Shift in Imitation Learning via Offline Data Without Great Coverage
Neural Information Processing Systems (**NeurIPS**), **2021**

ACADEMIC RESEARCH EXPERIENCE

Mila – Quebec AI Institute

Graduate Research Assistant

Sep 2024 - Present

Montréal, QC, Canada

- Focusing on developing data-efficient reinforcement learning and imitation learning algorithms, advised by Pablo Samuel Castro
- Co-led a project focusing on mitigating flaws in asymmetric actor-critic settings via making the critic less pessimistic (RLC 2025)

- Currently working on projects related to planning with geometric horizon models, as well as learning stronger behavior foundation models from offline data

Cornell University

Undergraduate/Graduate Researcher

Sep 2020 - Dec 2023

Ithaca, NY, USA

- Worked in Wen Sun's lab, where I assisted on projects focused on (1) joint representation learning for imitation learning in high-dimensional environments (e.g. Atari) and (2) model-based offline imitation learning in state-based, image-based and non-action-based graphics settings (NeurIPS 2021)
- Co-led a project focused on making Discriminator Actor-Critic more principled via gradient boosting methods (ICLR 2024)
- Other non-thesis projects included (1) self-predictive learning for RL in the image-based control context and (2) hybrid learning from preferences, both for RLHF-based and RL-free algorithms
- Assisted another student on a project focusing on using RL to finetune image generative models such as diffusion and consistency models

Mila – Quebec AI Institute

Research Collaborator

Apr 2021 - Mar 2022

Montréal, QC, Canada (remote)

- Reinforcement learning research for the LambdaZero project focusing on scaling drug discovery
- Looked into ways to improve exploration in GFlowNets using techniques such as epistemic uncertainty estimation, RND, and asymmetric self-play

INDUSTRY EXPERIENCE

Waymo

Research Intern

Jun 2026 - Sep 2026

London, UK

- Worked on RL for adversarial simulations with Luke Rowe, Asen Nachkov, Angad Singh, Omar Makhoul & Oleg Desheulin

Cohere

Research Intern

Sep 2025 - Feb 2026

Montréal, QC, Canada (remote, part-time)

- Built synthetic data pipelines for long-horizon environments with verifiable rewards. Specifically focusing on τ^2 -bench with the purpose of building strong SFT-trained models for post-training.
- Worked on RL post-training, improving GRPO performance on τ^2 -bench telecom
- Advised by Pat Verga, worked closely with Samuel Cahyawijaya, Florian Strub and Omar Darwiche Domingues

Cohere

Research Intern

Jun 2024 - Aug 2024

New York, NY, USA (remote)

- Worked on LLM post-training, with a focus on improving multistep mathematical reasoning
- Concurrently implemented SOTA preference learning algorithms and improved large language model training infrastructure as part of RL team needs
- Advised by Mohammad Gheshlaghi Azar & Olivier Pietquin

Apple

Research Intern

May 2022 - Sep 2022

Cupertino, CA, USA

- Fundamental RL research in the Apple MLR team
- Proposed using simple clustering of subtrajectory representations for offline option learning, advised by Walter Talbott. Resulting method was shown to be qualitatively much better than other methods at detecting behavioral differences across diverse offline datasets in the pixel-based DeepMind Control Suite, allowing for effective option learning and simpler offline RL
- Explored various different techniques for visual representation learning, including view-based and reconstruction-based methods
- Concurrently studied representation learning for on-policy RL with Riashat Islam & Devon Hjelm

Amazon Web Services

Jun 2021 - Aug 2021

Software Development Engineer Intern

Boston, MA, USA

- Worked on AWS Boost team, used Pandas and NumPy to (1) aggregate seller data across multiple time periods and (2) develop a performance metric based on available data to rank sellers on the platform
- Performance metric was aimed to be simple to compute, resulting in a linear model of different seller attributes that was a suitable ranking
- Integrated performance metric into a new page on the Boost web application with TypeScript

Cornell Cup Robotics

Oct 2020 - May 2021

Machine Learning Team Member

Ithaca, NY, USA

- Used Haystack API from DeepSet AI to develop a scalable Q/A system for an R2D2-like robot
- Offloaded all heavy-compute ML systems (~80% of compute) onto AWS to ease workload for the main machine

VMware Inc.

Jun 2020 - Aug 2020

Data Science Intern

Palo Alto, CA, USA (remote)

- Analyzed in-house device risk score model by comparing with ground-truth security scores across a diverse device dataset
- Constructed random forest models to determine which device features were most indicative of riskiness

SERVICES

Reviews

- **Conferences:** NeurIPS (2023, 2025), ICLR (2024, 2025, 2026), RLC (2025)

SKILLS

- **Languages:** Python, Java, OCaml, C++, C, $\text{\LaTeX}$
- **Libraries/Frameworks:** PyTorch, JAX, TensorFlow, NumPy, Pandas, SKLearn, PySpark, OpenCV, Git